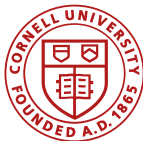


Detecting COVID-19 from Medical X-Ray Images

by Rayan Wali

Cornell University

December 2021



Motivation Behind Study

Problem Statement

When a medical image scan is taken by radiologists, they often have to decipher the image by classifying it as a single case.

The Problem

Given an X-ray medical image scan of a patient, decide whether the patient has COVID-19 or not.

AI as an Asset in Medical Imaging

- Reduces manual workload.
- Improves efficiency and results in precise decisions.
- Enables speedy delivery of results to patients.



Source: [The Computer Society](#)

Problem Statement (cont...)

We can encode this problem as a binary classification task, with the two classes being “Positive” and “Negative”.

Class #1: Positive

represents positive COVID-19 cases

Class #2: Negative

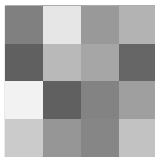
represents negative COVID-19 cases

Pixel Intensities

Each pixel of an image holds an intensity, and this property is used by the CNN to recognize the image as a whole. For an 256×256 dimensional image, there are in total 65,536 pixels. This pixel intensity array is then passed into the first layer of the CNN.



(a) Sample X-Ray Image



(b) Sample Grayscale Image

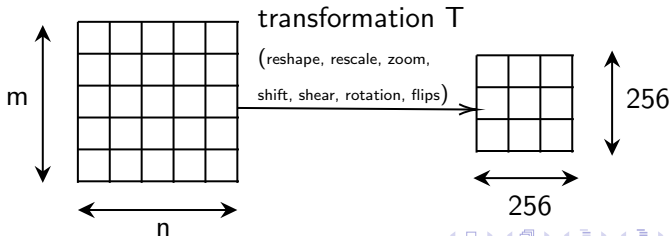
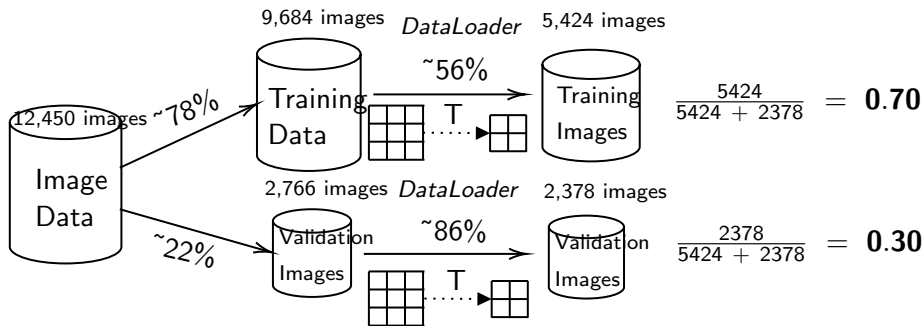
20	12	05	0	80
10	85	76	15	
22	510	22	78	
100	48	24	90	

(c) Pixel Matrix for (b)

- **Languages:** Python (Deep Learning Model), HTML (Webpage Structure), CSS (Webpage Styling).
- **Libraries/Packages:** Keras, TensorFlow (Data Preprocessing, Deep Learning Model Training + Testing), SkLearn (Data Preprocessing + Scientific Evaluation), Flask (Interactive Webpage).

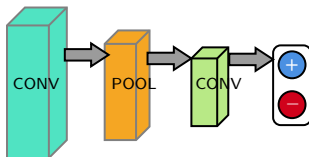
- ➊ **Preprocess** Input Data
- ➋ **Develop** Model Architecture
- ➌ **Train** the Model
- ➍ **Evaluate** the Model on Training and Validation Data
- ➎ **Deploy** Model for Web Application Use

Phase I: Preprocessing Data



Phase II: Developing the Model Architecture

VGG-16: A 16-layer CNN, with a mix of convolutional and pooling layers.



Model: "model"

Layer (type)	Output Shape	Param #
img_input (InputLayer)	[(None, 256, 256, 3)]	0
vgg16 (Functional)	(None, None, None, 512)	14714688
flatten (Flatten)	(None, 32768)	0
fc1 (Dense)	(None, 1024)	33555456
fc2 (Dense)	(None, 512)	524800
preds (Dense)	(None, 1)	513
Total params: 48,795,457		
Trainable params: 48,795,457		
Non-trainable params: 0		

Phase III: Training the Model

To train the model on the data, I used **stochastic gradient descent** (SGD). Stochastic gradient descent is an optimization technique that minimizes the following expression:

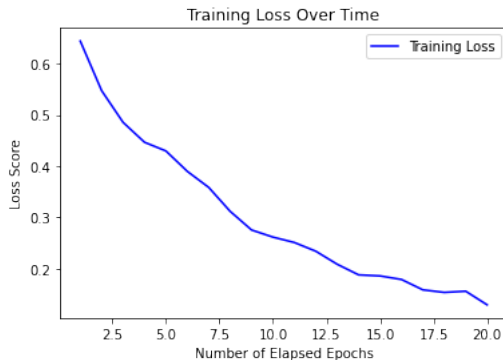
$$\min J(\theta) = \min \frac{1}{m} \sum_{i=1}^m \mathcal{L}(h_{\theta}(x_i), y_i) \quad (1)$$

HYPERPARAMETERS

- I trained my model for 20 epochs and achieved an 95% training accuracy in the end of the training process.
- For the gradient descent backpropagation phase, I used the binary cross-entropy loss function.
- Considering training time, accuracy, and loss, I found that $\alpha = 0.01$ was the optimal learning rate from the ones I experimented with.

Phase III: Visualizing the Training Process

In the following line plot, we observe (1) a continual decrease in the binary cross-entropy loss over time as the number of epochs that have elapsed grows. We also observe (2) that the rate of decline in the loss in the earlier epochs is greater than that in the later epochs.



Phase IV: Evaluating the Model

To evaluate my trained model, I considered three metrics outlined as follows:

- (1) **Accuracy Metric:** The number of correct instance classifications divided by the number of total instances classified.
- (2) **Loss Score:** May account for the confidence in the model's prediction in comparison to the true label. Is not limited to the simple evaluation strategy that the accuracy metric uses (0-1 loss function).
- (3) **F1 Score Metric:**

$$\text{F1 Score} = 2 \cdot \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}},$$

$$\text{where Precision} = \frac{TP}{TP + FP} \text{ and Recall} = \frac{TP}{TP + FN}.$$

Phase IV: Evaluating the Model (Accuracy)

[$\alpha = 0.01$, SGD Optimizer]

No. Epochs	Avg. Training Accuracy	Avg. Validation Accuracy
5	82%	81%
10	88%	86%
20	95%	96%

Phase IV: Evaluating the Model (F1 Score)

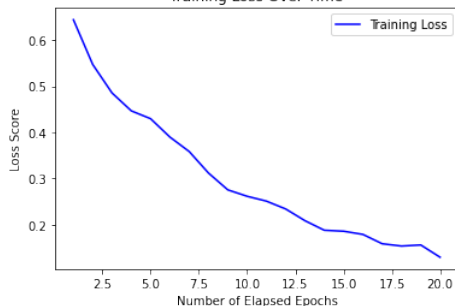
[$\alpha = 0.01$, SGD Optimizer]

No. Epochs	Training F1 Score	Validation F1 Score
5	0.8929	0.8734
10	0.9052	0.8805
20	0.9187	0.9125

Phase IV: Other Experiments (Optimizer) [$\alpha = 0.01$, 20 Epochs]

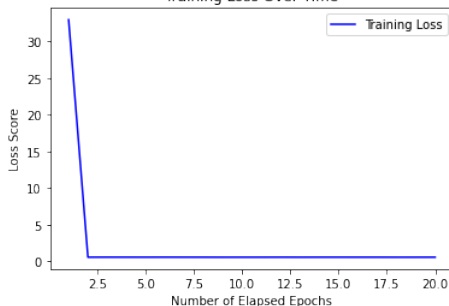
SGD

Training Loss Over Time



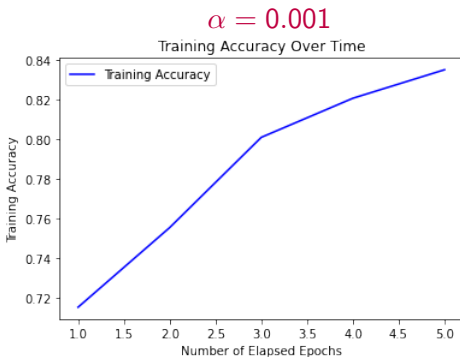
Adam

Training Loss Over Time



Phase IV: Other Experiments (Learning Rate α)

[SGD Optimizer, 5 Epochs]



Phase IV: Other Experiments – An Interesting Phenomenon (Learning Rate α) [SGD Optimizer, 5 Epochs]

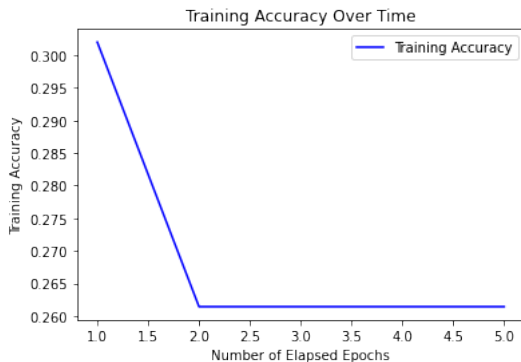
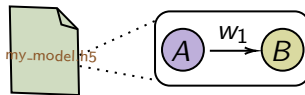


Figure: $\alpha = 0.1$ (when the learning rate is set too high, the training makes huge jumps; in this case, unexpected behavior occurs.)

Phase V: Deploying the Model

I built a web-application with Flask and HTML intended for use by radiologists who can input a patient's chest X-ray image and retrieve an output on whether the patient is COVID-19 positive or negative.

Model Compression (.h5 file):



COVID-19 Self-Test

Background

COVID-19 X-Ray Self-Testing Site



How Does The System Work?

This COVID-19 tool allows patients and radiologists to input 2-D chest X-Ray images. The system will process that image, and will return whether the patient whose 2-D chest X-ray image is inputted is COVID-positive or COVID-negative.

COVID-19 is known and studied to have a strong correlation with patient chest X-ray images, indicating that this system returns highly accurate predictions.

Instructions: Select a 2-D Chest X-Ray image and click the "Choose File" button below to perform the COVID test.

Choose File

No file chosen

Submit

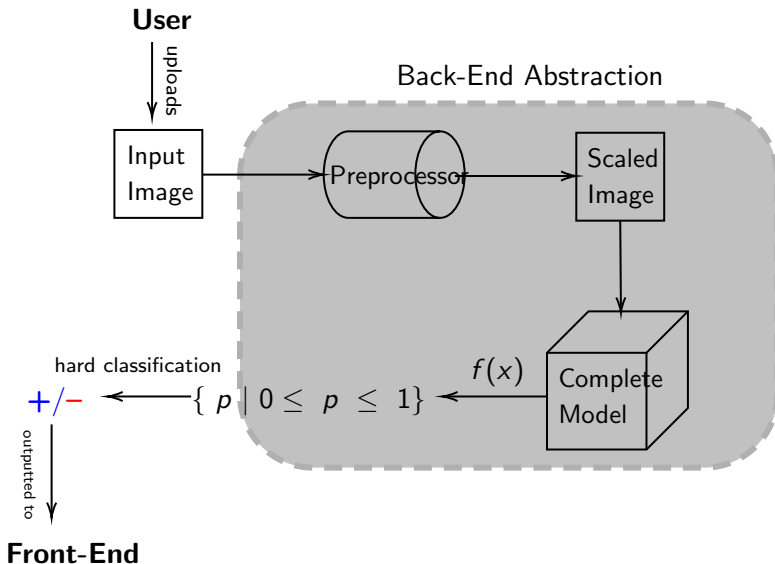
Normal-5.png



Model predicts: Negative

Learn more about COVID-19 [here](#).

End-to-End Picture



Potential Improvements & Optimizations

- 1 Perform **hyperparameter optimization** with random search or Bayesian optimization to obtain the optimal number of epochs and the optimizer type resulting in a lower overall loss.
- 2 Add more **data augmentation** to the training data in the preprocessing phase (to increase model generalizability).
- 3 Re-perform the training process on a different **model architecture**, e.g., on Microsoft's ResNet architecture.
- 4 Make the system **interactive**, i.e., the input of radiologists should be able to influence the output prediction.